

Wolt Test Assignment Report: Predicting number of couriers


MIKHAIL SILAEV

Introduction

►The problem of predicting time series is crucial for the future activity planning. The dataset given contains daily courier number combined with the weather data.

►Outlier-cleaned time series on this slide show visible correlations of “number of couriers” with temperature and precipitation. There is also a linear growing trend with time, probably due to the growing popularity of courier service.

►Is it possible to predict the number of couriers for next day or for several days ahead based on the historical data?

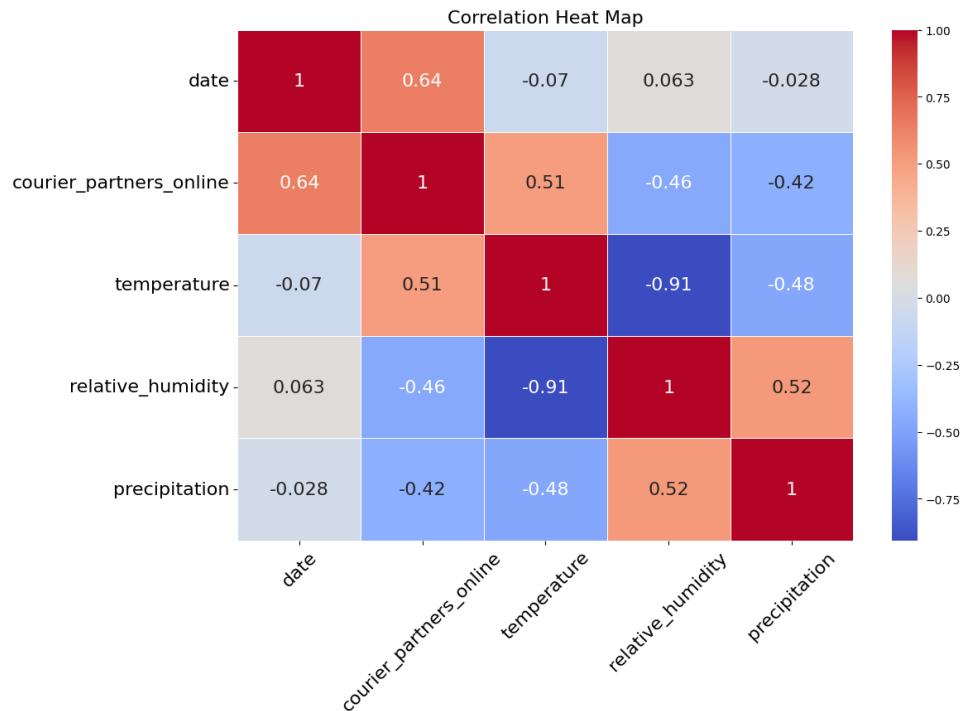


Exploratory Data Analysis

► Let us have a look at the correlation matrix first. Besides weather characteristics we add also 'day_of_week_i' dummy variable with $i = 0..6$. Here 'day_of_week_0' corresponds to Monday.

► Largest correlation for 'courier_partner_online' is with 'date', 'temperature', and then with 'precipitation'.

► Correlation with 'humidity' is also large. However, 'humidity' is almost deterministically correlated with 'temperature'. It should not add new prediction power.



Prediction Tasks & Feature Engineering

► Next-Day Task

Predicting the value of the y_{i+1} (representing 'courier_partners_online') for the next day $i+1$ based on the historical values available up to day i .

► Multiple-Day Task

Predicting the values of the variable $y_{i+1}, y_{i+2}, \dots, y_{i+n}$ for the next n days $i+1, \dots, i+n$ based on the historical values available up to day i .

► Feature selection

- Features consist historical variables 'courier_partner_online', 'temperature', 'precipitation', 'day_of_week_0', ... 'day_of_week_6'.
- Look-back window size: 'train_days' parameter controls number of preceding days to be used as features. We take typically 40 days to describe both small- and meso-time scales like season variability.
- With 10 feature and 40 preceding days, the total number of features is 400.

► Additional features

In the notebooks you can experiment adding more features such as 'date', periodic 'day_of_year', 'day_of_week'.

Models

Baseline - Linear Regression (LR) Model:


- **Parameters:** Equal to the number of features (400)

Long Short-Term Memory (LSTM) Model:


• Architecture:

- **Input:** (timesteps, features)
- **Layers:** 4 LSTM layers with varying units and dropout (30%, 10%, 20%, 30%)
- **Dense Layer:** Output layer with n units, which can be 1 or e.g. 20
- **Total params:** 138,271
- **Loss:** MSE, **Optimizer:** Adam


We adopt the [LSTM realization for stock price prediction](#)



Results: Next-Day prediction



Learning curves (see in /notebooks)



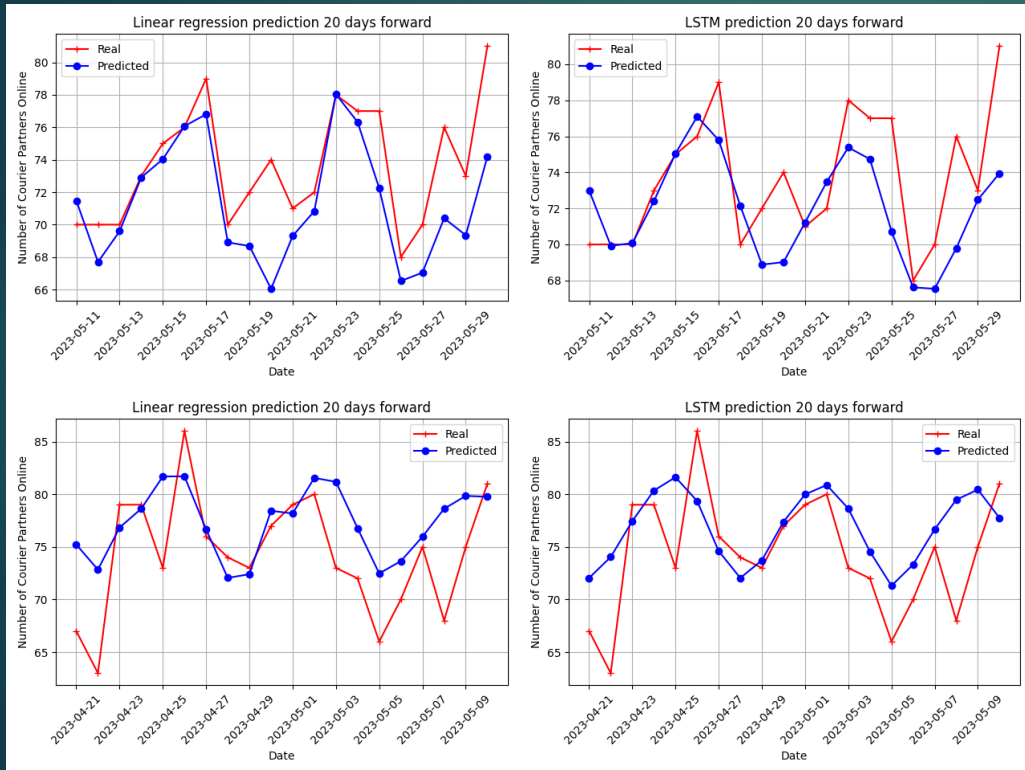
Residual
errors

Model	MAE	RMSE	SNR	R2
LSTM	2.92	3.93	25.61	0.35
LR	3.04	3.90	25.67	0.36

- Linear Regression works a bit better than LSTM
- Most important features: 'courier_partner_online', 'temperature'
- Optimal training stop is at 20 Epoch

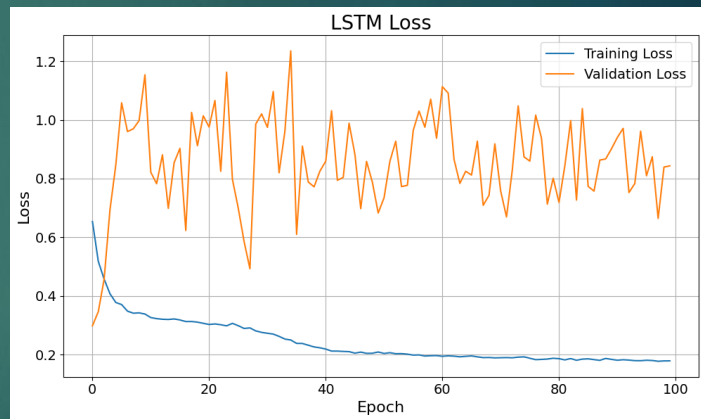
Results: Multiple-Day prediction for 20 ahead

Predictions for different 20-days terms



residual errors

Model	MAE	RMSE	SNR	R2
LSTM	3.43	4.32	24.95	0.29
LR	3.51	4.36	24.71	0.17



- LSTM performs better than LR
- However, LSTM overfits so that test and validation errors fluctuate. Metrics can vary from run to run

Conclusion

- ▶ Both LSTM and Linear Regression models predict courier numbers for the Next Day and for Multiple Days ahead based on the historical features.
- ▶ For both tasks metrics RMSE, SNR, and R-squared show decent performance (R-squared $\sim 0.2-0.3$). For the Next-Day task, both models perform similarly, while Linear Regression slightly outperforms for the Multiple-Day task.
- ▶ LSTM performance depends on hyperparameters like batch size and early stopping. For Multiple-Day task LSTM model overfits. Future improvements can focus on hyperparameter optimization and expanding the time series dataset to further enhance performance.